

3 Broken Mirror

3.1 What it is

Broken Mirror is a live mirror that shows you how an AI sees your face. You stand in front of the webcam, and the CLIP AI [1] — the same kind used by image search engines and photo apps — tries to describe you. It guesses your race, your gender, your age. It ranks hundreds of words by how well they “fit” your appearance, according to the patterns it learned from millions of photos on the internet. You see all of this in real time. The goal is not to get an accurate reading — the point is to see what the machine thinks it sees, and to notice where it is wrong, strange, unfair or funny.



Figure 3: Screenshot of the Broken Mirror tool UI.

3.2 UI walkthrough

Top-left — webcam panel. Your live camera feed. The Freeze button pauses it so you can discuss a specific moment without the image changing as you move.

Left sidebar — demographic guesses. The AI’s guesses about your race, gender and age, shown as percentage bars. In the screenshot, the AI is 49% sure the person is Indian, but also 19% sure they might be Middle Eastern and 16% sure they might be Latino/Hispanic. It’s 96% sure the person is female. Notice that these are guesses with confidence levels, not facts. A small movement, a different light, a different expression, and the numbers shift. That instability is part of what the tool is there to show.

Centre — the map. Imagine the AI organises all the words it knows into a giant mental landscape, where similar-meaning words sit close together and unrelated ones sit far apart. The map in the middle is a flattened

2D view of a tiny piece of that semantic landscape. Each small dot is a word. The purple star marked YOU is where your face lands in that landscape — which words the AI thinks you are closest to. In the screenshot, this person’s face has landed near “housekeeper”, “civil servant”, “clerk”, “receptionist”, “prisoner” and “prison officer”. The AI has decided these are the “closest” words to this face from a predefined list of words (see below).

Right sidebar — Category filters. The words on the map come from a published list of about 370 stereotype-loaded terms [2], grouped into ten themes: Appearance, Behavioural, Education and wealth, Criminal justice, Healthcare, Portrayal in media, Political, Religion, Occupation, and a Custom slot for words you add yourself. You can turn categories on or off to focus the map. In the screenshot (Figure 3), only Occupation is on, which is why the nearby words are all jobs.

Right sidebar — Add custom word. A text box where you can type any word, in any language, and see where it lands on the map. This is how you add your own words to the tool.

Right sidebar — Closest terms list. A ranked list of the words the AI considers most similar to your face right now, with a small number showing how close. In the screenshot, the top matches are housekeeper, prison officer, head teacher, solicitor, receptionist, prisoner, lawyer. This list is essentially “who the machine thinks you are”.

Session panel, language and display controls. A named session keeps your interactions grouped for later analysis. The top-right buttons change language (EN/ES), switch to light/dark mode, and adjust font size.

3.3 Proposed workshop dynamics

The tool is used in two parts: a guided demonstration led by the facilitator, and open play handed to participants, followed by a contribution phase.

3.3.1 Part 1 — Guided demonstration

The facilitator sits in front of the camera and walks everyone through the screen piece by piece. Start with the demographic guesses on the left: read the percentages out loud, then move your head slightly, change the light, smile, and show how the numbers jump around. The AI is not committing to one answer — it is spreading its guess across several, and that spread changes from second to second.

Then introduce the map and the word list with all categories active. Explain that “close to your face” does not mean “true about you” — it means these are words that tended to appear alongside faces like yours in the photos this AI learned from. The AI is not describing you. It is reflecting patterns it picked up from the internet.

Next, show the bias openly. Turn the categories on one by one and let the group see what comes up. Criminal justice shows words like “felon”, “prisoner”, “thug”. Healthcare shows words like “drug-seeking”, “insecure”, “disfigured”. Appearance shows words like “unattractive”, “sexy”, “groomed”.

Content warning — use before turning categories on

Warn the group before this step: *these categories contain offensive and derogatory terms on purpose* — the whole list was assembled by researchers precisely because these are the words that carry stereotypes. Seeing them projected onto a real face can be uncomfortable. Anyone who wants to step out should feel free to do so, and re-join when the demonstration moves on.

Then play with the AI. Put on glasses. Pull up a hood. Smile widely, then frown. Look down, look up. Try a “professional” posture, then a relaxed one. Bring in a second person. Hold up a book, a phone, a flower. Watch the demographic bars swing and the word list reshuffle. You will find that the AI is surprisingly easy to confuse

or push around, and the results are often absurd or funny. Lean into that. This is the moment where the group stops feeling watched by the machine and starts watching the machine back.

3.3.2 Part 2 — Open play

Turn off all categories except Occupation, which is the safest one to hand over — jobs rather than slurs. Pass the camera around. Let people try it on their own or in pairs. Try different faces, expressions, props, angles. Encourage talking: what did you expect? What surprised you? What felt wrong? The facilitator stays nearby but steps back.

3.3.3 Part 3 — Adding what's missing

Once everyone has played enough to have a sense of how the tool behaves, shift into the contribution phase. The Add custom word field is where this happens. The instructions are: “Looking at what the AI associates with you, and at what it clearly does not know about you — what words belong on this map that aren’t here? Add them and see what happens.”

Every word someone adds appears on the map immediately. Discuss a few of them out loud: where did this word land? Near which other words? And your face?