

4 Las Agencias

4.1 What it is

Las Agencias puts participants in the role of a communication team designing a real piece of public-facing material — a leaflet, a poster, a booklet for a local-government campaign. The campaign is concrete: “a new park in the city”, “a healthy-living initiative for the neighbourhood”, “a public library reopening”, “a neighbourhood assembly”. The participants’ job is to choose the photos that go into the graphic-design template.

The tool exists because every other audit in the project evaluates AI behaviour in the abstract. Broken Mirror asks “how does this AI see me?”. Fairness Arena asks “which AI ranks faces more fairly?”. These are powerful questions, but they are not how bias actually does harm in the real world. Harm happens when an AI is used *for something*, by someone with a task, under time pressure, with a template to fill. A communication agency asked to design a leaflet for a city park will run an image search, look at what comes back and pick what fits the layout. If the AI-assisted search returns only white citizens and no other people, that is what ends up on the leaflet — and that is what the city’s residents see. Las Agencias audits the AI in that exact role: as an assistant to a real communication task.

The campaign template is a normal booklet layout of the kind used by municipal communication offices (see Figure 4). It has several image slots of different sizes and proportions. The participants’ task is to fill those slots with images that (a) serve the campaign well and (b) fairly represent the community the campaign is meant to speak to.



Figure 4: Example booklet layout used in the Las Agencias tool.

4.2 Why this is a different kind of audit

Two things make this tool’s output meaningfully different from what Fairness Arena or Broken Mirror produce:

The judgment is situated. Participants are not asked “is this image biased?” in the abstract. They are asked “is this image good *for this campaign, for this community?*”. Context changes what counts as appropriate. An image that would be fine for a fashion magazine is wrong for a city-park leaflet; an image that would look stereotyped in isolation might be the only culturally-appropriate one in context. This kind of context-sensitive fairness judgment is what public-sector communication teams actually perform, and it is what the EU AI Act’s Fundamental Rights Impact Assessment framework expects deployers of AI systems to perform.

Erasure becomes visible. Who is not on the leaflet matters as much as who is. If the “healthy-living campaign” leaflet shows no older people, no disabled people, no one who looks poor, that is a representational harm even if every single image on the page is “good” in isolation. Las Agencias is the only tool in the project that surfaces

erasure at the page-composition level.

4.3 The image pool

The data-annotation pipeline that builds the image pool is documented separately under `intervisions_annotation` in the project organisation. The pool is constructed to satisfy three criteria simultaneously: *use-case validity* (each image plausibly fits at least one of the campaign briefs the tool runs), *diversity along perceived attributes* (gender, age, race, ability, visible markers of class), and *open licensing where possible*. Every image is documented with provenance and the perceived-attribute annotations used in analysis.

4.4 Two interaction modes

The tool supports two complementary interaction modes inside the same workshop format. The choice between them is workshop-specific and is set by the facilitator at the start of the session; both produce data that is comparable across models because both share the same set of campaign briefs and the same underlying image pool.

4.4.1 Mode A — Vote on pre-composed leaflets

The tool shows participants complete, laid-out booklets — the campaign template already populated with images chosen by different AI systems. For the same campaign brief, each booklet is filled with the top-N images in a search ranking obtained with a different AI. Participants then compare the booklets using best-worst scaling: “Which booklet is the best representation of this community? Which is the worst?”

This mode is simpler to run and keeps the workshop interaction fast and comparable. It also lets us re-use the same post-analysis as Fairness Arena — pairwise contrasts across the same AI variants, now measured at the page-composition level rather than at the individual-ranking level.

What it captures well. Participants evaluate the output as it would actually be seen — as a finished leaflet, not as an abstract ranking. Best-worst across three or four booklets per brief produces meaningful signal about which AI system produces better-composed material.

What it misses. Participants only see the top-N retrievals per model. If a particular group is absent from the top-N the booklet shows that absence, but participants have no way to inspect whether the group was present at rank N+1 or N+20 — they can only see what the AI placed at the top-8. This means subtle differences between models that only appear deeper in the ranking are invisible to this design. A model that is biased at rank 1–N but diverse at rank N+1–N+40 looks identical to a model that is biased throughout.

4.4.2 Mode B — Search-and-compose

The tool gives participants the campaign brief and the empty template, plus a search interface connected to a large image pool (thousands of images, not just the pre-selected pool). Participants type queries, browse results and drag chosen images into the template slots themselves. They decide what to search, what to reject and what to use.

The signal the tool captures is behavioural: which queries did they try? which results did they click? how far down the ranking did they scroll? how long did they spend before finding something usable? which images did they reject? The same task is done by different small groups across different CLIP variants behind the search — participants do not know which model they are using, they just know they are “doing the task” — and the tool

measures effort, success and satisfaction per model.

What it captures well. This is an ecologically valid audit of AI-assisted image search as it actually happens in practice. A model that forces participants to scroll through 40 results before finding a suitable image of an older person is worse than a model that puts such an image at rank 5, even if both eventually allow the task to be completed. Time and click-depth are meaningful efficiency metrics, and they are also fairness metrics when the effort gap is systematically larger for finding images of certain groups.

What it makes harder. The analysis is considerably more complex. With Mode A you have a clean vote per booklet. With Mode B you have behavioural traces per participant per model — click streams, scroll depths, query reformulations, final compositions. Meaningful metrics must be defined carefully: time-to-first-suitable-image, number-of-queries-needed, depth-of-accepted-images, rejection patterns. The ground-truth relevance question — which images in the pool are “genuinely suitable” for the brief — remains open and is supported by the perceived-attribute annotations from the pipeline above.

4.4.3 Hybrid use across workshops

The two modes are also designed to be chained across workshops: the compositions that participants build in Mode B in one workshop can become the booklets that other participants (or an online cohort) vote on in Mode A in a later session. This recovers the simplicity of Mode A’s voting while grounding the booklets in real human composition decisions rather than pure top- k retrieval, and it doubles the signal density of each Mode B participant’s contribution.